

TS

March 25, 2018

```
In [1]: import numpy as np
        from math import ceil, log, log10, sqrt, exp
        import matplotlib.pyplot as plt
        from scipy.stats import beta
        import pandas as pd
        import dill

In [2]: # Saving & Loading Variables
        filename = 'globalsave.pkl'
        # dill.load_session(filename)

In [3]: # for reproducibility
        np.random.seed(1234)

In [4]: # UCB Implementation given horizon (time steps), #replications, True arm means & Type
        # For plotting average % Optimal arm pulls
        def UCB(horizon, replications, arms_prob, ucctype, optimalpulls):

            optimal_arm = 0

            optimal_arm_pulls_per_round = np.zeros([horizon, replications]) # Stores % optimal
            regret_per_round = np.zeros([horizon, replications]) # Stores regret for every time

            for r in range(replications):
                arm_means = [0]*len(arms_prob) # Initializing arm means & pulls to 0
                arm_pulls = [0]*len(arms_prob)
                t = 0
                #initially playing each arm once
                for i in range(len(arms_prob)):
                    arm_pulls[i] += 1
                    temp = np.random.binomial(1, arms_prob[i])
                    arm_means[i] += (temp - arm_means[i])/arm_pulls[i] # Updating arm means es
                    if optimalpulls == 'Avg':
                        optimal_arm_pulls_per_round[t][r] = arm_pulls[optimal_arm]*100.0/(t+1)
                    else:
                        if i == optimal_arm: # Incrementing % optimal arm pulls if current arm
                            optimal_arm_pulls_per_round[t][r] += 1
                        regret_per_round[t][r] = (arms_prob[optimal_arm] - arms_prob[i]) # Storing
```

```

t+=1

while t < horizon:
    #Picking arm according to UCB algorithm or UCB' algorithm
    if ucctype == 1:
        UCBestimate = arm_means + np.sqrt(2*np.log(t)/arm_pulls)
        arm_selected = np.argmax(UCBestimate)
    else:
        UCBestimate = arm_means + np.sqrt(2*np.log(horizon)/arm_pulls)
        arm_selected = np.argmax(UCBestimate)

    arm_pulls[arm_selected] += 1
    temp = np.random.binomial(1, arms_prob[arm_selected]) # Updating arm means
    arm_means[arm_selected] += (temp - arm_means[arm_selected]) / arm_pulls[arm_selected]
    regret_per_round[t][r] = (arms_prob[optimal_arm] - arms_prob[arm_selected])
    if optimalpulls == 'Avg':
        optimal_arm_pulls_per_round[t][r] = arm_pulls[optimal_arm]*100.0/(t+1)
    else:
        if arm_selected == optimal_arm: # Incrementing % optimal arm pulls if
            optimal_arm_pulls_per_round[t][r] += 1
    t+=1

# Calculating Mean and Standard Error for % optimal arm pulls
optimal_arm_means_stderr = np.zeros([horizon,2]) # Store % optimal arm means & stderr
optimal_arm_means_stderr[:,0] = np.mean(optimal_arm_pulls_per_round,axis=1)
optimal_arm_means_stderr[:,1] = (np.std(optimal_arm_pulls_per_round, axis=1)/sqrt(repetitions))
if optimalpulls == 'Avg':
    optimal_arm_percentage = sum(optimal_arm_means_stderr[:,0])/horizon
    optimal_arm_pulls_sum = np.mean(optimal_arm_pulls_per_round,axis=1)
else:
    optimal_arm_percentage = sum(optimal_arm_means_stderr[:,0])/horizon*100
    optimal_arm_pulls_sum = np.cumsum(optimal_arm_means_stderr[:,0])/horizon*100
print("\nTotal Optimal arm pulls :",sum(optimal_arm_means_stderr[:,0]),'and percentage')

# Calculating Mean and Standard Error for commulative regret
regret_means_stderr = np.zeros([horizon,2]) # Store regret means & stderr in (horizon, repetitions)
regret_means_stderr[:,0] = np.mean(regret_per_round,axis=1)
regret_means_stderr[:,1] = (np.std(regret_per_round, axis=1)/sqrt(repetitions))
total_regret = sum(regret_means_stderr[:,0])
regret_per_round_sum = np.cumsum(regret_means_stderr[:,0])
print("Total Regret :",total_regret)

return regret_per_round_sum,regret_means_stderr, optimal_arm_pulls_sum,optimal_arm_percentage

```

```

In [5]: # Implementation Thompson Sampling
def TS(horizon,replications,arms_prob,alpha,beta,optimalpulls):
    optimal_arm = 0
    optimal_arm_pulls_per_round = np.zeros([horizon,replications])
    regret_per_round = np.zeros([horizon,replications])
    savepoints = (0,1000,5000,9999)
    success_ret = np.zeros([len(savepoints),len(arms_prob)])
    faliure_ret = np.zeros([len(savepoints),len(arms_prob)])

    for r in range(replications):
        arm_pulls = [0]*len(arms_prob)
        success = np.array(alpha)
        failure = np.array(beta)
        t = 0
        s = 0
        while t < horizon:
            if t in savepoints and r == replications-1:
                success_ret[s] = success
                faliure_ret[s] = failure
                s+=1

            #Picking arm according to Posterior distribution
            sample_means = [0]*len(arms_prob)
            for i in range(len(arms_prob)):
                sample_means[i] = np.random.beta(success[i],failure[i])
            arm_selected = np.argmax(sample_means)

            arm_pulls[arm_selected] += 1
            temp = np.random.binomial(1, arms_prob[arm_selected])
            success[arm_selected] += temp
            failure[arm_selected] += 1 - temp

            if optimalpulls == 'Avg':
                optimal_arm_pulls_per_round[t][r] = arm_pulls[optimal_arm]*100.0/(t+1)
            else:
                if arm_selected == optimal_arm: # Incrementing % optimal arm pulls if
                    optimal_arm_pulls_per_round[t][r] += 1

            regret_per_round[t][r] = (arms_prob[optimal_arm] - arms_prob[arm_selected])

            t+=1

    # Calculating Mean and Standard Error for % optimal arm pulls
    optimal_arm_means_stderr = np.zeros([horizon,2])
    optimal_arm_means_stderr[:,0] = np.mean(optimal_arm_pulls_per_round,axis=1)
    optimal_arm_means_stderr[:,1] = (np.std(optimal_arm_pulls_per_round, axis=1)/sqrt(

```

```

if optimalpulls == 'Avg':
    optimal_arm_percentage = sum(optimal_arm_means_stderr[:,0])/horizon
    optimal_arm_pulls_sum = np.mean(optimal_arm_pulls_per_round,axis=1)
else:
    optimal_arm_percentage = sum(optimal_arm_means_stderr[:,0])/horizon*100
    optimal_arm_pulls_sum = np.cumsum(optimal_arm_means_stderr[:,0])/horizon*100

print("\nTotal Optimal arm pulls :",sum(optimal_arm_means_stderr[:,0]),'and percent

# Calculating Mean and Standard Error for commulative regret
regret_means_stderr = np.zeros([horizon,2])
regret_means_stderr[:,0] = np.mean(regret_per_round,axis=1)
regret_means_stderr[:,1] = (np.std(regret_per_round, axis=1)/sqrt(replications))
total_regret = sum(regret_means_stderr[:,0])
regret_per_round_sum = np.cumsum(regret_means_stderr[:,0])
print("Total Regret :",total_regret)

return success_ret,faliure_ret,regret_per_round_sum,regret_means_stderr, optimal_a

```

In [6]: # Calculating Lower bound Regret

```

def calculate_lower_bound():
    arms_prob = [[0.5, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4], [0.5, 0.48, 0.48,

n = 10000
gap_dependent_regret = np.zeros([3,n])
gap_independent_regret = np.zeros([3,n])
for i in range(3):
    k = len(arms_prob[i])
    for j in range(1,n):
        dep_reg = 0
        for g in range(len(arms_prob[i])):
            gap = arms_prob[i][0] - arms_prob[i][g]
            if gap > 0 and dep_reg >= 0:
                dep_reg += log(j*gap**2)/(8*gap) #n*gap * exp(-4*j*gap**2/(k-1))/4
            if dep_reg < 0:
                dep_reg = 0
        gap_dependent_regret[i][j] = dep_reg

        indep_reg = sqrt((k-1)*j/8)*exp(-1/2)/4
        gap_independent_regret[i][j] = indep_reg
return gap_dependent_regret,gap_independent_regret

```

```
gap_dependent_regret,gap_independent_regret = calculate_lower_bound()
```

```

print('gap_dependent_regret')
print(gap_dependent_regret)

print('\ngap_independent_regret')
print(gap_independent_regret)

gap_dependent_regret
[[ 0.          0.          0.          ...,  51.80478909  51.80591437
  51.80703954]
 [ 0.          0.          0.          ...,  77.96218028  77.96780669
  77.97343253]
 [ 0.          0.          0.          ...,   5.13966203   5.13973496
   5.13980789]]

gap_independent_regret
[[ 0.          0.16083073  0.227449   ...,  16.0806602  16.08146446
  16.08226867]
 [ 0.          0.16083073  0.227449   ...,  16.0806602  16.08146446
  16.08226867]
 [ 0.          0.07581633  0.10722049 ...,   7.58049592   7.58087505
   7.58125416]]

In [7]: # Plotting % Cumulative Optimal Arm Pulls Vs Time steps with error bars
def plotCumOptimalArmPulls(horizon,optimal_arm_means_stderr,optimal_arm_pulls_sum,prob
    x = np.arange(horizon)
    ind = [i for i in range(0,horizon,step)]

    for i in range(m_len):
        plt.errorbar(x[ind],optimal_arm_pulls_sum[i,ind], optimal_arm_means_stderr[i,i
            linestyle='-', marker='x', capsize=4, capthick=1.5, elinewidth=1.5)
    plt.xlabel('Steps')
    plt.ylabel('% Cumulative Optimal Arm Pulls')
    plt.legend(['UCB', "TS with prior mean=0.5", "TS with prior mean=0.2"],loc=0,frameon=
    plt.title(arms_prob[problem])
    plt.xlim((0,10000))
    plt.ylim((0,100))
    plt.savefig('CummOptimalArmPulls_'+str(problem)+optimalpulls+'.png',dpi=300)
    plt.show()

    print("optimal_arm_stderr")
    print(optimal_arm_means_stderr[:, [500,2000,5000,8000,9500],1])

In [8]: # Plotting % Average Optimal Arm Pulls Vs Time steps with error bars
def plotAvgOptimalArmPulls(horizon,optimal_arm_means_stderr,optimal_arm_pulls_sum,prob
    x = np.arange(horizon)
    ind = [i for i in range(0,horizon,step)]

```

```

for i in range(m_len):
    plt.errorbar(x[ind], optimal_arm_pulls_sum[i, ind], optimal_arm_means_stderr[i, ind],
                 linestyle='-', marker='x', capsize=4, capthick=1.5, elinewidth=1.5)
plt.xlabel('Steps')
plt.ylabel('% Average Optimal Arm Pulls')
plt.legend(['UCB', "TS with prior mean=0.5", "TS with prior mean=0.2"], loc=0, frameon=True)
plt.title(arms_prob[problem])
plt.xlim((0, 10000))
plt.ylim((0, 100))
plt.savefig('AvgOptimalArmPulls_'+str(problem)+optimalpulls+'.png', dpi=300)
plt.show()

print("optimal_arm_means_stderr")
print(optimal_arm_means_stderr[:, [500, 2000, 5000, 8000, 9500], 1])

In [9]: # Plotting Cumulative Regret Vs Time steps with error bars
def plotCummRegret(horizon, regret_means_stderr, regret_per_round_sum, problem, step):
    labels = ["Gap-dependent lower bound", "Gap-independent lower bound", 'UCB', "TS with"]
    x = np.arange(horizon)
    ind = [i for i in range(0, horizon, step)]

    plt.plot(x[ind], gap_dependent_regret[problem][ind], label = labels[0])
    plt.plot(x[ind], gap_independent_regret[problem][ind], label = labels[1])

    for i in range(m_len):
        plt.errorbar(x[ind], regret_per_round_sum[i, ind], regret_means_stderr[i, ind, 1],
                     linestyle='-', marker='x', capsize=4, capthick=1.5, elinewidth=1.5)

    plt.xlabel('Steps')
    plt.ylabel('Cumulative Regret')
    plt.legend(loc=0, frameon=False)
    plt.title(arms_prob[problem])
    plt.xlim((0, 10000))
    # plt.ylim((0, 80))
    plt.savefig('CumulativeRegret_'+str(problem)+optimalpulls+'.png', dpi=300)
    plt.show()

    print("regret_means_stderr")
    print(regret_means_stderr[:, [500, 2000, 5000, 8000, 9500], 1])
    print("regret_per_round_sum")
    print(regret_per_round_sum[:, [500, 2000, 5000, 8000, 9500]])
    print("gap_dependent_regret")
    print(gap_dependent_regret[problem, [500, 2000, 5000, 8000, 9500]])
    print("gap_independent_regret")
    print(gap_independent_regret[problem, [500, 2000, 5000, 8000, 9500]])

In [10]: def plot_arm_distribution(alpha_values, beta_values):

```

```

savepoints = [1,1000,5000,10000]
s = 0
for p in range(len(alpha_values)):

    fig, ax = plt.subplots(figsize=(10, 6))
    m = alpha_values[p] / (alpha_values[p] + beta_values[p])
    print("Mean at time t = "+str(savepoints[s])+" is: ",m)

    x = np.linspace(0, 1, 1000)[1:-1]
    i = 1
    for a, b in zip(alpha_values[p], beta_values[p]):
        plt.plot(x, beta.pdf(x,a,b),
                 label=r'Arm %d : $\alpha$=%d,\ \beta=%d$' % (i, a, b))
        i+=1

    plt.ylim(0, 20)
    plt.xticks(np.arange(0, 1.1,0.1))
    plt.xlabel(r'$ \theta $')
    plt.ylabel(r'$p(\theta|\alpha,\beta)$')
    plt.title('Problem '+str(problem+1)+' at t = '+str(savepoints[s]))
    s+=1
    plt.legend(loc=0)
    plt.savefig('Arm_dist_t_'+str(savepoints[s-1])+'_'+str(problem+1)+'_'+str(savepoints[s]))
    plt.show()

```

```

In [11]: # For Printing table
from IPython.display import HTML, display

```

```

def tableIt(data):
    print(pd.DataFrame(data))

```

```

In [12]: horizon = 10000
replications = 100

```

```

arms_prob = [[0.5, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4], [0.5, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48]]
types = ['UCB', 'TS M0.5', 'TS M0.2']
m_len = len(types)
success = [1,1]
failure = [1,4]
optimalpulls = 'Avg'

for problem in range(3): # Repeating for 3 problems
    optimal_arm_pulls_sum = np.zeros([m_len,horizon]) # Storing variables returned by
    regret_per_round_sum = np.zeros([m_len,horizon])
    optimal_arm_means_stderr = np.zeros([m_len,horizon,2])
    regret_means_stderr = np.zeros([m_len,horizon,2])
    optimal_arm_percentage = np.zeros([m_len])

```

```

total_regret = np.zeros([m_len])
success_ret = np.zeros([2,4,len(arms_prob[problem])])
failure_ret = np.zeros([2,4,len(arms_prob[problem])])

regret_per_round_sum[0,:],regret_means_stderr[0,:,:], optimal_arm_pulls_sum[0,:],
success_ret[0,:,:],failure_ret[0,:,:],regret_per_round_sum[1,:],regret_means_stderr[1,:,:],
success_ret[1,:,:],failure_ret[1,:,:],regret_per_round_sum[2,:],regret_means_stderr[2,:,:]

step = 500
print("\n")
print("optimal_arm_percentage")
tableIt(optimal_arm_percentage)
print("\n")
print("total_regret")
tableIt(total_regret)

# Calling function to plot % Average Optimal Arm Pulls & Commulative regret with
plotAvgOptimalArmPulls(horizon,optimal_arm_means_stderr,optimal_arm_pulls_sum,problem)
plotCummRegret(horizon,regret_means_stderr,regret_per_round_sum,problem,step)
plot_arm_distribution(success_ret[0],failure_ret[0])
plot_arm_distribution(success_ret[1],failure_ret[1])

```

Total Optimal arm pulls : 283612.675268 and percentage is : 28.3612675268
Total Regret : 593.993

Total Optimal arm pulls : 652499.168408 and percentage is : 65.2499168408
Total Regret : 163.781

Total Optimal arm pulls : 683731.231325 and percentage is : 68.3731231325
Total Regret : 145.473

optimal_arm_percentage

```

0
0 28.361268
1 65.249917
2 68.373123

```


total_regret

```

0
0 593.993
1 163.781

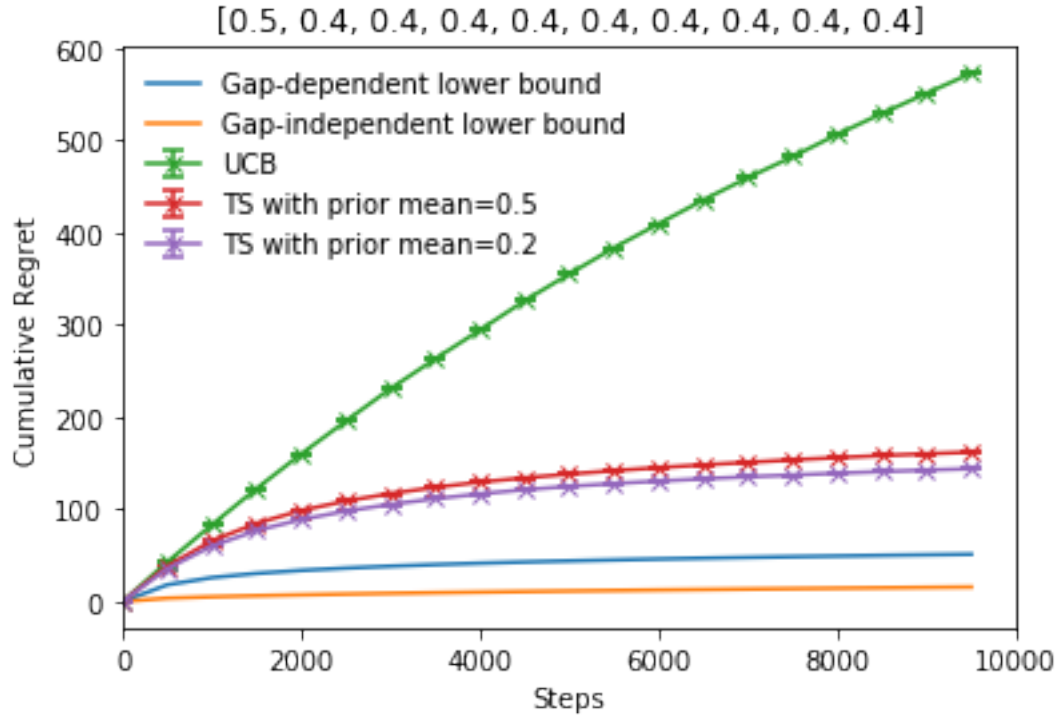
```


2 145.473



optimal_arm_means_stderr


```
[[ 0.37932051  0.43692854  0.44865863  0.45474807  0.435826 ]
 [ 1.37482323  1.32860494  0.6824189   0.44764855  0.38093131]
 [ 2.01556859  1.74196148  1.1765315   0.72305736  0.6041692 ]]
```




```

regret_means_stderr
[[ 0.00324962  0.00392301  0.00497494  0.00491833  0.00462493]
 [ 0.00448999  0.00384187  0.00255147  0.00099499  0.00099499]
 [ 0.00495076  0.00420833  0.00237487  0.00217945  0.00255147]]
regret_per_round_sum
[[ 42.736  160.02  355.174  506.605  572.967]
 [ 38.318  98.747  138.166  156.021  161.991]
 [ 35.768  88.902  124.882  139.08  144.003]]
gap_dependent_regret
[ 18.10617651  33.70198808  44.01025881  49.29779964  51.23111503]
gap_independent_regret
[ 3.59628442  7.19256883  11.37244987  14.38513767  15.67584034]
Mean at time t = 1 is: [ 0.5  0.5  0.5  0.5  0.5  0.5  0.5  0.5  0.5  0.5]

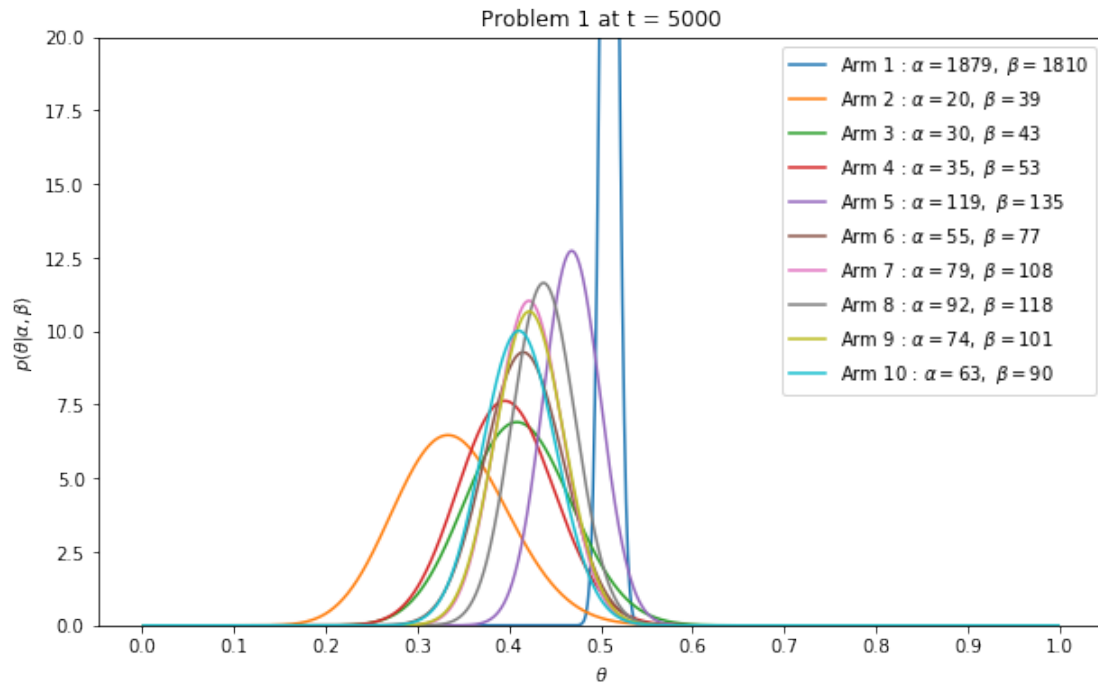
```




Mean at time $t = 1000$ is: [0.52199413 0.34693878 0.30769231 0.38636364 0.45833333 0.445
0.44444444 0.46551724 0.39344262 0.44444444]




Mean at time $t = 5000$ is: [0.50935213 0.33898305 0.4109589 0.39772727 0.46850394 0.41666667
0.42245989 0.43809524 0.42285714 0.41176471]




Mean at time $t = 10000$ is: [0.50435757 0.34883721 0.37 0.38655462 0.45081967 0.40384615
0.42205323 0.41666667 0.41803279 0.42801556]




Mean at time $t = 1$ is: [0.2 0.2 0.2 0.2 0.2 0.2 0.2 0.2 0.2 0.2]




Mean at time $t = 1000$ is: [0.49002217 0.3902439 0.31034483 0.1 0.32 0.348
0.41578947 0.39735099 0.09090909 0.1]



Mean at time $t = 5000$ is: [0.48611867 0.36419753 0.26415094 0.42937853 0.40888889 0.385
0.40718563 0.40677966 0.17647059 0.25]



Mean at time t = 10000 is: [0.49659624 0.36144578 0.26415094 0.40756303 0.41698842 0.39040895522 0.4140625 0.19047619 0.24324324]



Total Optimal arm pulls : 120911.777501 and percentage is : 12.0911777501
Total Regret : 173.3864

Total Optimal arm pulls : 186662.859336 and percentage is : 18.6662859336
Total Regret : 151.8254

Total Optimal arm pulls : 184985.5294 and percentage is : 18.49855294
Total Regret : 151.7754

optimal_arm_percentage

```

0
0 12.091178
1 18.666286
2 18.498553


```

total_regret

```

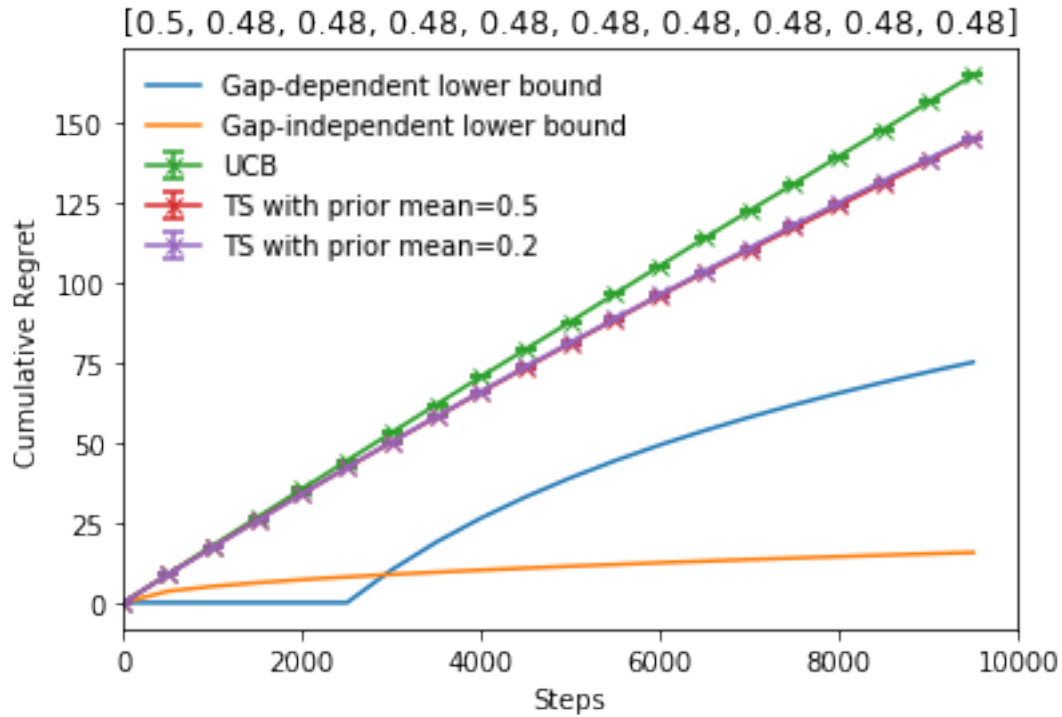
0
0 173.3864
1 151.8254
2 151.7754

```



optimal_arm_means_stderr

```
[[ 0.28318001  0.28274657  0.29367604  0.30803299  0.30607703]
 [ 0.85373802  1.1549887   1.30079714  1.3876149   1.41370241]
 [ 1.26060158  1.35384415  1.46942818  1.54023541  1.56892188]]
```



regret_means_stderr

```
[[ 0.00069397  0.00043589  0.00064992  0.00071414  0.00069397]
 [ 0.00071414  0.00084167  0.00081462  0.00092499  0.00094742]
 [ 0.00073321  0.0008      0.00084167  0.00091652  0.00096   ]]
```

regret_per_round_sum

```
[[ 8.952   35.5116  87.7602 139.3692 164.7034]
 [ 8.807   34.0506  80.9434 123.9802 144.9186]
 [ 8.7878  33.9774  81.2388 124.8678 145.241  ]]
```


gap_dependent_regret

```
[ 0.          0.          38.98952891  65.42723305  75.09381   ]
```


gap_independent_regret

```
[ 3.59628442  7.19256883 11.37244987 14.38513767 15.67584034]
```


Mean at time t = 1 is: [0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5]




Mean at time $t = 1000$ is: [0.51485149 0.35483871 0.524 0.44594595 0.515625 0.304
0.46666667 0.36666667 0.46590909 0.49]




Mean at time $t = 5000$ is: [0.51424361 0.38666667 0.49730216 0.44444444 0.47184987 0.3170
0.47193878 0.31707317 0.47493404 0.47188264]




Mean at time $t = 10000$ is: [0.51450163 0.41558442 0.49598624 0.43220339 0.47201493 0.49
0.47684391 0.38372093 0.4745098 0.47758887]



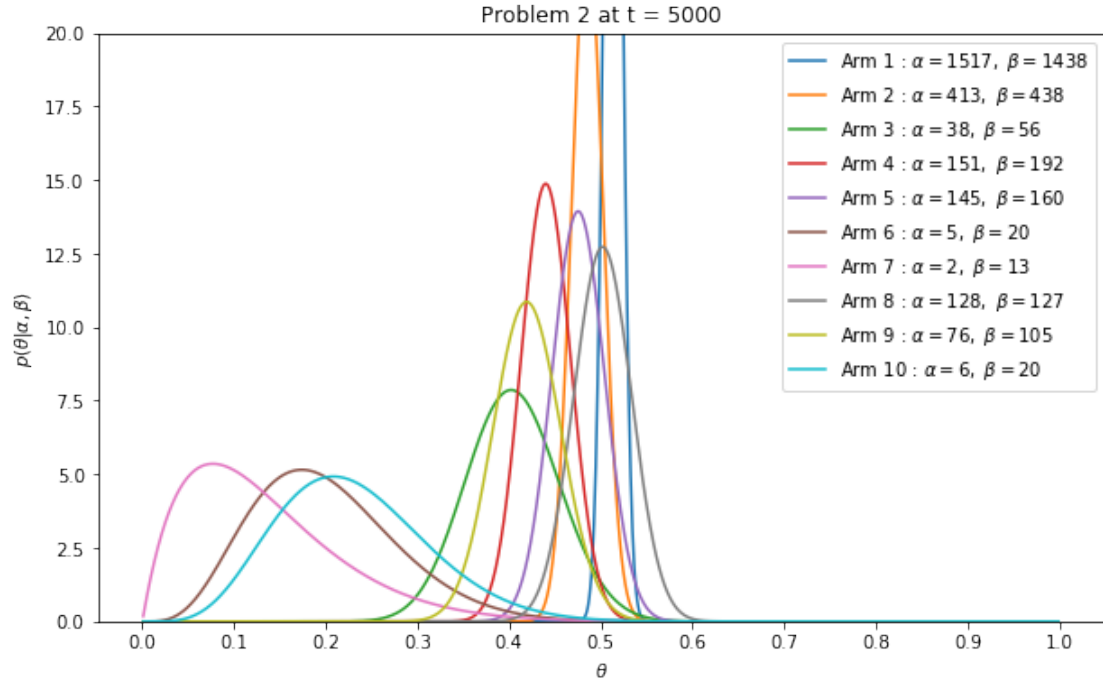
Mean at time t = 1 is: [0.2 0.2 0.2 0.2 0.2 0.2 0.2 0.2 0.2 0.2]




Mean at time $t = 1000$ is: [0.50588235 0.52631579 0.25 0.47849462 0.31578947 0.235
0.18181818 0.41666667 0.32142857 0.1]



Mean at time $t = 5000$ is: [0.51336717 0.4853114 0.40425532 0.44023324 0.47540984 0.2
0.13333333 0.50196078 0.4198895 0.23076923]



Mean at time t = 10000 is: [0.49947424 0.48360656 0.38970588 0.45631068 0.47182176 0.300
0.11764706 0.48771121 0.45733788 0.28571429]



Total Optimal arm pulls : 934693.661168 and percentage is : 93.4693661168
 Total Regret : 83.607

Total Optimal arm pulls : 980563.54202 and percentage is : 98.056354202
 Total Regret : 16.274

Total Optimal arm pulls : 984437.531967 and percentage is : 98.4437531967
 Total Regret : 12.754

optimal_arm_percentage


```

0
0 93.469366
1 98.056354
2 98.443753
    
```

total_regret


```

0
0 83.607
1 16.274
2 12.754
    
```



optimal_arm_means_stderr

```
[[ 0.30812797  0.11263889  0.05128719  0.03391501  0.02926298]
 [ 0.32196661  0.0802825   0.03251039  0.02044813  0.01726094]
 [ 0.29373779  0.07279044  0.02945255  0.01821061  0.01522703]]
```



regret_means_stderr

```
[[ 0.00964365  0.00397995  0.00397995  0.00298496  0.00298496]
 [ 0.00298496  0.00298496  0.          0.          0.          ]
 [ 0.00298496  0.          0.          0.          0.          ]]
```

regret_per_round_sum

```
[[ 32.377  55.085  71.444  79.962  82.712]
 [ 10.293  13.056  14.944  15.859  16.159]
 [  7.919   9.769  11.291  12.24   12.671]]
```


gap_dependent_regret

```
[ 2.95549269  3.96633233  4.63446099  4.97717197  5.10247945]
```


gap_independent_regret

```
[ 1.69530473  3.39060946  5.36102428  6.78121893  7.38966201]
```


Mean at time t = 1 is: [0.5 0.5 0.5]




Mean at time $t = 1000$ is: [0.49693252 0.13333333 0.15384615]




Mean at time $t = 5000$ is: [0.49728316 0.11764706 0.15]



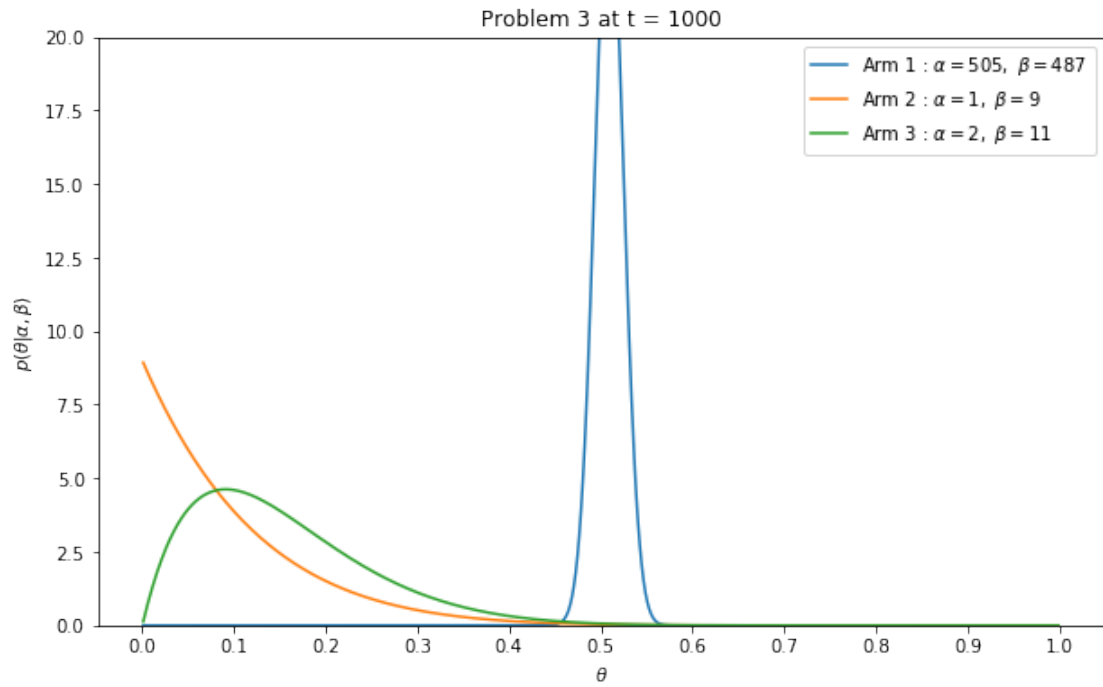
Mean at time $t = 10000$ is: [0.50060205 0.11111111 0.14285714]




Mean at time $t = 1$ is: [0.2 0.2 0.2]




Mean at time $t = 1000$ is: [0.50907258 0.1 0.15384615]



Mean at time t = 5000 is: [0.50391802 0.11764706 0.14285714]



Mean at time $t = 10000$ is: [0.50365915 0.11111111 0.14285714]



```
In [13]: horizon = 10000
         replications = 100
```

```
arms_prob = [[0.5, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4], [0.5, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48], [0.5, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48]]
types = ['UCB', 'TS M0.5', 'TS M0.2']
m_len = len(types)
success = [1,1]
failure = [1,4]
optimalpulls = 'Cum'
```

```
for problem in range(3): # Repeating for 3 problems
    optimal_arm_pulls_sum = np.zeros([m_len,horizon]) # Storing variables returned by
    regret_per_round_sum = np.zeros([m_len,horizon])
    optimal_arm_means_stderr = np.zeros([m_len,horizon,2])
    regret_means_stderr = np.zeros([m_len,horizon,2])
    optimal_arm_percentage = np.zeros([m_len])
    total_regret = np.zeros([m_len])
    success_ret = np.zeros([2,4,len(arms_prob[problem])])
    failure_ret = np.zeros([2,4,len(arms_prob[problem])])
```

```
regret_per_round_sum[0,:],regret_means_stderr[0,:,:], optimal_arm_pulls_sum[0,:],
```

```

success_ret[0,:,:],failure_ret[0,:,:],regret_per_round_sum[1,:],regret_means_stderr[1,:],
success_ret[1,:,:],failure_ret[1,:,:],regret_per_round_sum[2,:],regret_means_stderr[2,:]

```

```

step = 500
print("\n")
print("optimal_arm_percentage")
tableIt(optimal_arm_percentage)
print("\n")
print("total_regret")
tableIt(total_regret)

```

```

# Calling function to plot % Optimal Arm Pulls Vs Time steps with error bars
plotCumOptimalArmPulls(horizon,optimal_arm_means_stderr,optimal_arm_pulls_sum,prob)
plot_arm_distribution(success_ret[0],failure_ret[0])
plot_arm_distribution(success_ret[1],failure_ret[1])

```

```

Total Optimal arm pulls : 4085.01 and percentage is : 40.8501
Total Regret : 591.499

```

```

Total Optimal arm pulls : 8373.46 and percentage is : 83.7346
Total Regret : 162.654

```

```

Total Optimal arm pulls : 8530.6 and percentage is : 85.306
Total Regret : 146.94

```

```

optimal_arm_percentage

```

```

0
0 40.8501
1 83.7346
2 85.3060

```

```


total_regret

```

```

0
0 591.499
1 162.654
2 146.940

```




optimal_arm_stderr

[[0.03249615 0.04439595 0.04918333 0.04828043 0.04950758]


[0.04737088 0.04439595 0.02179449 0.02179449 0.01705872]

[0.04877499 0.03756328 0.02712932 0.00994987 0.01959592]]


Mean at time t = 1 is: [0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5]




Mean at time $t = 1000$ is: [0.49578059 0.44615385 0.32142857 0.30232558 0.28571429 0.392
0.26666667 0.35135135 0.25 0.2]




Mean at time $t = 5000$ is: [0.49002101 0.43413174 0.38709677 0.3375 0.28571429 0.382
0.27272727 0.43362832 0.21875 0.38043478]




Mean at time $t = 10000$ is: [0.49940688 0.42465753 0.44290657 0.32941176 0.28 0.37
0.25531915 0.43089431 0.24390244 0.4021164]




Mean at time $t = 1$ is: [0.2 0.2 0.2 0.2 0.2 0.2 0.2 0.2 0.2 0.2]




Mean at time $t = 1000$ is: [0.3877551 0.3125 0.36507937 0.27906977 0.25714286 0.1
0.43595041 0.40340909 0.37349398 0.23076923]



Mean at time $t = 5000$ is: [0.49939976 0.39181287 0.37606838 0.2962963 0.36507937 0.436
0.42495479 0.40692641 0.37007874 0.20689655]



Mean at time t = 10000 is: [0.49506435 0.38709677 0.3961039 0.28070175 0.42965779 0.44
0.42576029 0.39583333 0.37956204 0.20689655]



Total Optimal arm pulls : 1351.15 and percentage is : 13.5115
Total Regret : 172.977

Total Optimal arm pulls : 2266.69 and percentage is : 22.6669
Total Regret : 154.6662


Total Optimal arm pulls : 2225.47 and percentage is : 22.2547
Total Regret : 155.4906

optimal_arm_percentage

```
0
0 13.5115
1 22.6669
2 22.2547
```

total_regret


```
0
0 172.9770
1 154.6662
2 155.4906
```




optimal_arm_stderr

```
[[ 0.03          0.03249615  0.03666061  0.0346987   0.03570714]
 [ 0.03249615  0.04          0.04702127  0.04439595  0.04624932]
 [ 0.03          0.03128898  0.04208325  0.04828043  0.04582576]]
```


Mean at time $t = 1$ is: [0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5]




Mean at time $t = 1000$ is: [0.41666667 0.48979592 0.484375 0.14285714 0.48192771 0.494
0.125 0.54397394 0.50485437 0.42307692]




Mean at time t = 5000 is: [0.53886616 0.48070175 0.45495495 0.35849057 0.46808511 0.5
0.42352941 0.51771117 0.5106383 0.40196078]




Mean at time $t = 10000$ is: [0.51920744 0.47954545 0.46179402 0.41836735 0.48235294 0.49
0.46927374 0.49976202 0.47916667 0.42857143]




Mean at time $t = 1$ is: [0.2 0.2 0.2 0.2 0.2 0.2 0.2 0.2 0.2 0.2]




Mean at time $t = 1000$ is: [0.47852761 0.33333333 0.53527981 0.32258065 0.46788991 0.181
0.42424242 0.53465347 0.41025641 0.48421053]



Mean at time $t = 5000$ is: [0.49255751 0.31428571 0.48929664 0.49479941 0.47147147 0.217
0.42307692 0.49968691 0.37383178 0.47044335]



Mean at time t = 10000 is: [0.49607623 0.33333333 0.47994467 0.48519219 0.48195576 0.19
0.45983936 0.49320794 0.38095238 0.4754717]



Total Optimal arm pulls : 9749.23 and percentage is : 97.4923
Total Regret : 84.45

Total Optimal arm pulls : 9952.5 and percentage is : 99.525
Total Regret : 16.055


Total Optimal arm pulls : 9963.59 and percentage is : 99.6359
Total Regret : 12.261

optimal_arm_percentage

```
0
0 97.4923
1 99.5250
2 99.6359
```

total_regret


```
0
0 84.450
1 16.055
2 12.261
```




optimal_arm_stderr

```
[[ 0.01705872  0.01959592  0.          0.00994987  0.014       ]
 [ 0.00994987  0.          0.          0.          0.          ]
 [ 0.00994987  0.00994987  0.          0.          0.          ]]
```


Mean at time t = 1 is: [0.5 0.5 0.5]




Mean at time t = 1000 is: [0.47983454 0.24 0.14285714]




Mean at time t = 5000 is: [0.5 0.25 0.11111111]




Mean at time $t = 10000$ is: [0.49270697 0.24444444 0.10526316]




Mean at time $t = 1$ is: [0.2 0.2 0.2]




Mean at time $t = 1000$ is: [0.509684 0.125 0.16666667]



Mean at time $t = 5000$ is: [0.50914205 0.11764706 0.14285714]



Mean at time t = 10000 is: [0.51298246 0.11111111 0.14285714]




```

In [14]: # Saving variables
         dill.dump_session(filename)

In [15]: def get_arm_playing_prob(success,failure,trials=10000):

         sample_means = np.zeros([trials,len(success)])
         arm_played_times = np.zeros(trials)
         arm_paying_prob = np.zeros(len(success))
         for i in range(len(success)):
             sample_means[:,i] = np.random.beta(success[i],failure[i],trials)

         for j in range(trials):
             arm_played_times[j] = np.argmax(sample_means[j,:])

         for i in range(len(success)):
             arm_paying_prob[i] = (arm_played_times == i).sum()/trials

         return arm_paying_prob

In [31]: # Plotting Arm Playing Probability Vs Time steps for TS
         def plot_arms_playing_prob(horizon,arm_play_Prob,problem):
             x = np.arange(horizon)
             for i in range(len(arm_play_Prob[1,:])):
                 plt.plot(x,arm_play_Prob[:,i])
             plt.xlabel('Steps')
             plt.ylabel('Arm Playing Probability per step')
             if len(arm_play_Prob[1,:]) == 3:
                 plt.legend(["Arm 1","Arm 2","Arm 3"],loc="best",frameon=False)
             else:
                 plt.legend(["Arm 1","Arm 2","Arm 3","Arm 4","Arm 5","Arm 6","Arm 7","Arm 8","Arm 9"],loc="best",frameon=False)
             if problem % 2 == 0:
                 plt.title("Problem 3: TS with prior mean = 0.2")
             else:
                 plt.title("Problem 3: TS with prior mean = 0.5")

             plt.xlim((0,10000))
             plt.ylim((0,1))
             plt.savefig('Arm_play_prob_'+str(problem)+'.png',dpi=300)
             plt.show()

In [17]: # Implementation Thompson Sampling for plotting arm playing probability at each time step
         def TSForPlottingArmProb(horizon,replications,arms_prob,alpha,beta,optimalpulls,problem):
             optimal_arm = 0

             # optimal_arm_pulls_per_round = np.zeros([horizon,replications])
             # regret_per_round = np.zeros([horizon,replications])
             arm_playing_prob = np.zeros([horizon,len(arms_prob)])

```

```

savepoints = (0,1000,5000,9999)
success_ret = np.zeros([len(savepoints),len(arms_prob)])
faliure_ret = np.zeros([len(savepoints),len(arms_prob)])

for r in range(replications):
    arm_pulls = [0]*len(arms_prob)
    success = np.array(alpha)
    failure = np.array(beta)
    t = 0
    s = 0
    while t < horizon:
        if t in savepoints and r == replications-1:
            success_ret[s] = success
            faliure_ret[s] = failure
            s+=1

        #Picking arm according to Posterior distribution
        sample_means = [0]*len(arms_prob)
        for i in range(len(arms_prob)):
            sample_means[i] = np.random.beta(success[i],failure[i])
        arm_selected = np.argmax(sample_means)

        arm_playing_prob[t] = get_arm_playing_prob(success,failure)

        arm_pulls[arm_selected] += 1
        temp = np.random.binomial(1, arms_prob[arm_selected])
        success[arm_selected] += temp
        failure[arm_selected] += 1 - temp

        t+=1

    # plot_arms_playing_prob(horizon,arm_playing_prob,problem)
    return arm_playing_prob



```

```

In [18]: arm_playing_prob1 = TSForPlottingArmProb(10000,1,[0.5, 0.2, 0.1],[1,1,1],[1,1,1],'Avg
plot_arms_playing_prob(10000,arm_playing_prob1,1)


arm_playing_prob2 = TSForPlottingArmProb(10000,1,[0.5, 0.2, 0.1],[1,1,1],[4,4,4],'Avg
plot_arms_playing_prob(10000,arm_playing_prob2,2)


```



```
In [19]: arm_playing_prob3 = TSForPlottingArmProb(10000,1,[0.5, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4],
plot_arms_playing_prob(10000,arm_playing_prob3,3)
```



```
arm_playing_prob4 = TSForPlottingArmProb(10000,1,[0.5, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4, 0.4],
plot_arms_playing_prob(10000,arm_playing_prob4,4)
```






```
In [20]: arm_playing_prob5 = TSForPlottingArmProb(10000,1,[0.5, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48],
plot_arms_playing_prob(10000,arm_playing_prob5,5)


arm_playing_prob6 = TSForPlottingArmProb(10000,1,[0.5, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48, 0.48],
plot_arms_playing_prob(10000,arm_playing_prob6,6)
```




```
In [21]: dill.dump_session(filename)
```


```
In [23]: plot_arms_playing_prob(10000,arm_playing_prob3,3)  
         plot_arms_playing_prob(10000,arm_playing_prob4,4)
```







```
In [25]: plot_arms_playing_prob(10000,arm_playing_prob5,5)
         plot_arms_playing_prob(10000,arm_playing_prob6,6)
```





```
In [30]: plot_arms_playing_prob(10000,arm_playing_prob1,1)  
         plot_arms_playing_prob(10000,arm_playing_prob2,2)
```



```
In [32]: plot_arms_playing_prob(10000,arm_playing_prob1,1)
         plot_arms_playing_prob(10000,arm_playing_prob2,2)
```

